The AI Alpha Paradox.

Deterministic Multi-AI Signal Construction for Institutional Portfolio Governance

Ivan Alessandro Ongaro
Founder – AI Market Neutral
IOOX Srl – Milan, Italy

https://aimarketneutral.com

March 2026

Working Paper

Draft for discussion purposes. Not for redistribution.

Table of Contents

Abstract
1. Introduction
2. The AI Alpha Paradox
2.1 From Prediction to Structured Exploration
2.2 Analytical Pluralism as a Design Principle
3. Governance-First Signal Design
3.1 The Institutional Priority Hierarchy
3.2 Inverting the Hierarchy
3.3 Determinism as an Institutional Requirement
3.4 Signal-as-a-Service: Separation of Signal and Execution
3.5 Architecture Overview
4. Multi-Agent Analytical Layer
4.1 Agent Architecture and Independence
4.2 Signal Agent 1: Growth and Momentum (Claude Sonnet 4.5)
4.3 Signal Agent 2: Risk and Downside Fragility (Mistral Large 2512)
4.4 Signal Agent 3: Neutral Factor-Based Analysis (Gemini 3 Flash Preview)
4.5 Signal Agent 4: Contrarian Mean-Reversion (Grok 4)
4.6 Orchestrator: Final Arbitration (Claude Haiku 4.5)
4.7 Portfolio Critic: Two-Level Validation
5. Deterministic Signal Construction Engine
5.1 Rule Hierarchy and Execution Logic
5.2 Equal-Weight Design Philosophy
5.3 Correlation and Diversification Controls
5.4 Event Risk Management
5.5 Reproducibility and Immutability
6. Beta-Neutral Signal Framework
6.1 Investment Universe Selection
6.2 Long/Short Structure
6.3 Beta Neutrality and Market Exposure Management
6.4 Periodic Signal Generation
7. Auditability and Decision Integrity
7.1 Decision Record Architecture
7.2 Cryptographic Immutability
7.3 Transparency for Institutional Stakeholders
8. Implications for Institutional AI Deployment
8.1 The Complementarity Thesis

8.2 Implications for Regulatory Frameworks .....
8.3 The Signal-as-a-Service Model .....
8.4 Scalability Across Investment Strategies .....
8.5 The Human Element .....
9. Limitations and Future Research .....
9.1 Absence of Empirical Evidence .....
9.2 Agent Design and Model Selection .....
9.3 Regime Dependence .....
9.4 Integration with Macroeconomic Signals .....
9.5 Transaction Cost Modeling .....
10. Conclusion .....
Author Information .....
References .....
Disclaimer .....

Abstract

Recent advances in generative artificial intelligence have stimulated significant interest in the application of AI systems to financial markets. However, most implementations rely on opaque predictive models operating with limited structural constraints, raising concerns regarding governance, explainability, and risk control in institutional portfolio management.

This paper proposes an alternative framework: a deterministic signal construction architecture assisted by multiple independent AI agents, in which artificial intelligence is subordinated to a rule-based risk engine governing signal formation. The system is designed to operate as a Signal-as-a-Service platform, delivering structured, auditable portfolio signals to institutional subscribers rather than executing trades or managing capital directly.

The proposed system employs multiple AI analytical perspectives—each implemented through a distinct large language model with a specialized analytical orientation—to generate candidate securities within a defined investment universe. These proposals are subsequently filtered, constrained, and assembled by a deterministic signal construction engine designed to enforce strict exposure, correlation, liquidity, and event-risk controls.

The architecture aims to address the fundamental tension between AI-driven idea generation and institutional requirements for transparency, auditability, and deterministic risk governance. We present the conceptual design of the framework and discuss its application to the generation of beta-neutral long/short signals within the NASDAQ-100 universe.

Keywords: *multi-agent AI, deterministic signal construction, beta-neutral strategies, Signal-as-a-Service, governance-first architecture, institutional risk management, systematic equity strategies*

1. Introduction

Artificial intelligence has rapidly become a central topic in quantitative finance. Numerous firms and academic studies have explored machine learning models for predicting asset returns, identifying market regimes, or optimizing portfolio allocation. The proliferation of large language models and generative AI systems has further accelerated this trend, with an increasing number of market participants seeking to integrate AI capabilities into their investment processes.

Despite these advances, institutional adoption of AI-driven investment strategies remains markedly limited. While proprietary trading desks and quantitative hedge funds have experimented with various machine learning techniques, the broader institutional asset management industry—pension funds, sovereign wealth funds, insurance companies, and regulated investment managers—has been considerably more cautious. The gap between technological capability and institutional adoption represents one of the most significant challenges in modern quantitative finance.

One major obstacle is the lack of deterministic governance structures around AI-generated signals. Most AI-driven strategies rely on predictive models whose internal logic is difficult to interpret and whose outputs may vary significantly across training iterations or model updates. This opacity poses fundamental challenges for fiduciary investors operating under strict regulatory and compliance frameworks that demand explainability, reproducibility, and clear chains of accountability.

The concerns surrounding AI adoption in institutional portfolio management can be organized into several categories. First, there is the problem of limited explainability: when an AI system recommends a position, it is frequently impossible to articulate in precise terms why that recommendation was made, making it difficult to satisfy regulatory requirements for investment rationale documentation. Second, decision processes may be unstable, as model retraining on slightly different data windows can produce materially different outputs without clear justification. Third, the absence of deterministic audit trails means that the exact sequence of decisions leading to a signal composition cannot always be reconstructed after the fact. Finally, governance and compliance frameworks designed for human decision-making processes do not readily accommodate the stochastic nature of AI outputs.

In practice, portfolio managers operating within institutional mandates require systems where risk governance precedes signal generation—not the reverse. The traditional approach in quantitative finance has been to define the risk budget, exposure limits, and portfolio constraints before any alpha signal is considered. AI systems that generate unconstrained recommendations and then attempt to fit them within risk parameters after the fact represent an inversion of this established hierarchy that is fundamentally incompatible with institutional practice.

This paper introduces a design approach that restores the traditional governance hierarchy while incorporating the analytical capabilities of modern AI systems. In the proposed architecture, AI serves as an analytical layer—a source of structured exploration across diverse perspectives—while signal construction remains strictly deterministic. The framework separates the creative, pattern-recognition function of AI from the disciplined, rule-based function of signal formation, allowing each component to operate within its domain of greatest effectiveness. Critically, the system is designed as a Signal-as-a-Service platform: it delivers

structured, fully auditable portfolio signals to institutional subscribers via API, without executing trades or managing capital directly. This separation of signal generation from execution preserves the governance integrity of the signal while allowing each subscriber to implement the signal within their own risk and execution frameworks.

The architecture described in this paper is implemented within the AI Market Neutral signal platform, available at https://aimarketneutral.com.

The remainder of this paper is organized as follows. Section 2 examines the structural paradox inherent in applying AI to financial markets. Section 3 presents the governance-first design philosophy. Section 4 describes the multi-agent analytical layer, including the specific large language models deployed and their analytical orientations. Section 5 specifies the deterministic signal construction engine. Sections 6 through 8 discuss the beta-neutral framework, auditability mechanisms, and broader implications for institutional AI deployment, respectively. Section 9 addresses limitations and future research directions, and Section 10 concludes.

2. The AI Alpha Paradox

The application of AI to financial markets faces a structural paradox that is rarely addressed explicitly in the literature but is well understood among practitioners. This paradox has profound implications for how AI systems should be designed and deployed within investment processes, and its recognition is essential for building architectures that deliver sustainable value.

On one hand, AI systems excel at discovering patterns across large, heterogeneous information sets. Modern language models and machine learning architectures can process vast quantities of textual, numerical, and relational data, identifying subtle regularities that may elude traditional quantitative screens. The capacity to synthesize information across different modalities—financial statements, market microstructure data, news sentiment, sector dynamics, management commentary—represents a genuine analytical advance over single-factor or purely statistical approaches.

On the other hand, financial markets are adaptive systems populated by sophisticated participants who rapidly exploit and thereby erode any discovered predictive signal. Unlike physical systems governed by stable natural laws, markets are reflexive: the act of trading on a signal alters the conditions that generated it. This adaptive quality means that predictive signals tend to decay rapidly once they become widely known or exploited, a phenomenon extensively documented in the academic literature on factor crowding, alpha decay, and the diminishing marginal returns to quantitative research (McLean & Pontiff, 2016).

2.1 From Prediction to Structured Exploration

Consequently, the long-term value of AI in signal generation may not lie primarily in the direct prediction of asset returns. Single-model prediction architectures face a fundamental sustainability challenge: any model that successfully predicts returns will, once deployed at scale, alter market dynamics in ways that undermine its own predictive validity. This is not a technical limitation that can be overcome with better algorithms or more data; it is a structural feature of competitive markets that reflects the efficient market hypothesis in its semi-strong form.

The architecture proposed in this paper reframes the role of AI in signal construction. Rather than relying on a single predictive model to generate alpha signals, the system employs multiple AI agents—each implemented through a distinct large language model—with deliberately different analytical biases. Each agent approaches the investment universe from a different analytical perspective, producing diverse candidate sets of securities. This design draws conceptually from ensemble methods in machine learning (Dietterich, 2000), but with a critical distinction: the agents are not combined through statistical weighting to produce a consensus prediction. Instead, their outputs serve as input to a deterministic construction engine that applies its own independent logic, entirely divorced from the analytical process.

2.2 Analytical Pluralism as a Design Principle

The concept of analytical pluralism is central to the proposed architecture. By deliberately constructing agents with different analytical orientations—growth momentum, downside fragility, factor-based analysis, and mean reversion—the system generates a broad set of

candidate securities that reflect multiple views of the market. No single agent's perspective dominates the signal, and no agent has the authority to determine the final composition.

This approach acknowledges a fundamental epistemic limitation: no single analytical framework consistently captures all relevant aspects of financial market dynamics. Growth momentum signals may be effective in trending markets but fail in mean-reverting environments. Fundamental analysis may identify long-term value but miss short-term catalysts. Risk-based approaches may preserve capital in drawdowns but underperform in risk-on environments. By maintaining multiple perspectives simultaneously, the system avoids excessive dependence on any single market narrative and creates a structurally diversified analytical foundation.

Importantly, the diversity of analytical perspectives is not a hedge against model error in the traditional sense. It is a structural feature designed to ensure that the candidate set presented to the deterministic engine is rich, varied, and unlikely to be dominated by a single factor exposure or market regime assumption. The value of diversity lies not in error cancellation but in the breadth of the opportunity set it generates.

3. Governance-First Signal Design

Institutional portfolio construction has historically been governed by a clear hierarchy of priorities. At the apex of this hierarchy sit the risk budget and governance constraints that define the boundaries within which investment decisions must be made. Below these structural constraints, specific investment signals and alpha sources are evaluated and implemented. This ordering is not arbitrary; it reflects decades of institutional experience with the consequences of allowing investment ideas to precede risk discipline (Grinold & Kahn, 1999).

3.1 The Institutional Priority Hierarchy

The standard institutional priority framework encompasses several layers. Risk budgets define the maximum acceptable level of portfolio volatility, tracking error, or drawdown. Exposure limits constrain the portfolio's sensitivity to specific factors, sectors, or regions. Liquidity constraints ensure that positions can be entered and exited within acceptable timeframes and market impact parameters. Correlation controls limit the degree to which portfolio positions are exposed to common risk factors (Meucci, 2005). Event risk management addresses the potential for discrete events—earnings announcements, regulatory actions, corporate restructurings—to produce outsized position-level returns.

In many AI-driven investment systems, these constraints are applied after signals are generated. The AI model produces a set of predicted returns or position recommendations, and a post-hoc optimization layer attempts to reshape these recommendations to satisfy risk constraints. This approach suffers from a fundamental architectural weakness: the optimization layer must compromise between the AI's recommendations and the risk constraints, often producing outputs that are unstable across rebalancing periods, sensitive to minor input changes, and difficult to interpret or explain to institutional stakeholders.

3.2 Inverting the Hierarchy

In the proposed architecture, the hierarchy is explicitly reversed. Deterministic governance rules are defined first, establishing the complete set of constraints that any signal must satisfy before any consideration of specific securities or analytical assessments. AI agents operate only within the boundaries defined by these governance rules, generating candidate securities that are pre-filtered for basic eligibility criteria. Signal construction is then executed by a rule-based engine that mechanically assembles the output according to predetermined logic.

This inversion has several important consequences. The risk profile of the signal is determined entirely by the governance rules, not by the AI agents. Changes to the signal's risk characteristics can only occur through explicit modifications to the governance framework, creating a clear governance audit trail that connects every change to a deliberate rule modification. The AI agents cannot, by design, override or circumvent the risk framework, regardless of the strength of their analytical assessments. This separation provides a natural division of responsibilities: the governance framework is owned by the risk management function, while the AI agents are managed by the research function.

3.3 Determinism as an Institutional Requirement

The term "deterministic" is used deliberately throughout this paper to emphasize a specific property of the signal construction process: given the same set of inputs (candidate securities and governance rules), the engine will always produce the same signal. This property is essential for institutional adoption because it enables complete reproducibility of the signal generation process, facilitates regulatory examination and audit, allows independent verification of construction logic, and eliminates the possibility of hidden discretionary decisions within the construction process.

Determinism does not imply rigidity. The governance rules themselves can be updated periodically to reflect changes in the investment environment, risk appetite, or market structure. However, at any given point in time, the mapping from inputs to signal is fully specified and reproducible. This distinction between the static determinism of the construction process and the dynamic evolution of the governance framework is an important feature of the architecture, allowing it to adapt over time while maintaining transparency at every point.

3.4 Signal-as-a-Service: Separation of Signal and Execution

A defining characteristic of the proposed architecture is its explicit separation of signal generation from trade execution. The system does not manage capital, execute trades, or interact with brokerage infrastructure. Instead, it delivers structured portfolio signals—specifying long and short constituents, weights, sector allocations, and risk characteristics—to institutional subscribers via API.

This Signal-as-a-Service model has several important implications. It allows the signal provider to focus exclusively on analytical quality and governance integrity, without the regulatory and operational complexity of asset management. It allows each subscriber to implement the signal within their own execution framework, applying their own transaction cost models, market impact estimates, and timing preferences. And it creates a clean separation between the intellectual property embedded in the signal construction process and the execution infrastructure required to act on it. For institutional subscribers, this separation ensures that adopting the signal does not require ceding control over their execution process or custody arrangements.

3.5 Architecture Overview

The following diagram illustrates the end-to-end signal construction pipeline, from multi-agent analysis through deterministic construction to API delivery.

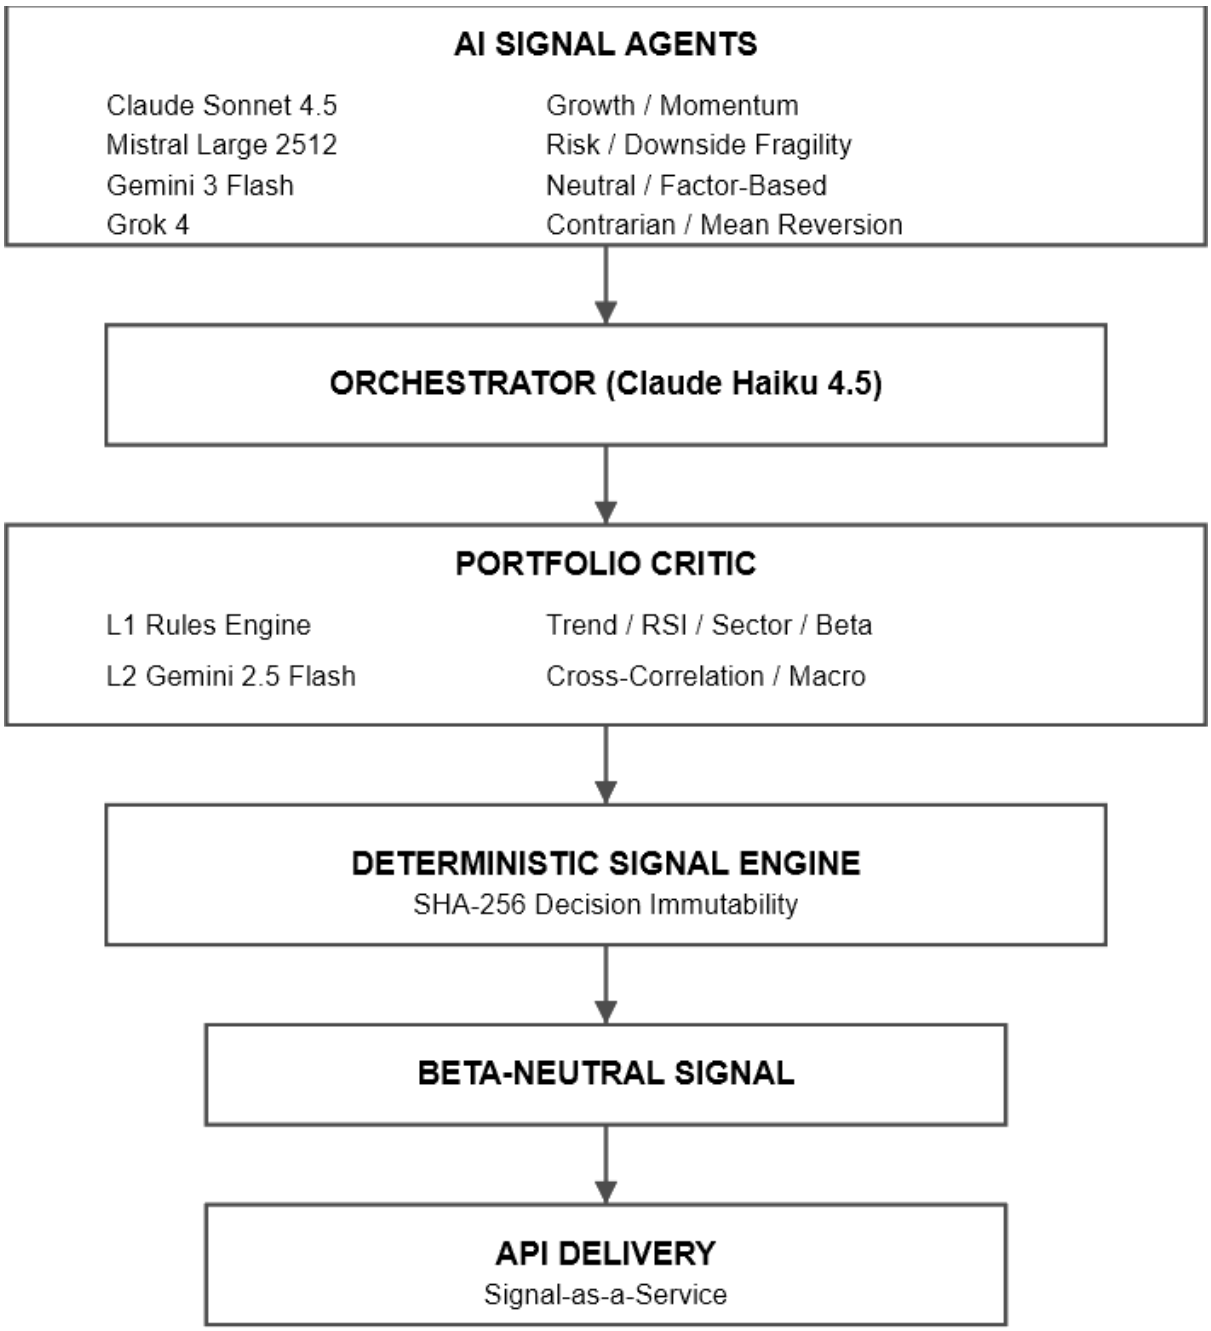

Figure 1. End-to-end signal construction architecture.
No human override at any stage. All decisions are deterministic and SHA-256 hashed.

4. Multi-Agent Analytical Layer

The system employs multiple independent AI agents acting as analytical perspectives rather than predictive authorities. This distinction is critical to understanding the architecture's design philosophy. The agents do not claim to predict future returns with precision; they provide structured analytical assessments that inform the candidate selection process, offering the deterministic engine a diverse and well-characterized set of securities from which to construct the signal.

A distinctive feature of this architecture is that each agent is implemented through a different large language model, selected for the specific analytical characteristics that align with the agent's designated orientation. This model diversity provides an additional layer of structural independence: not only do the agents evaluate the universe from different analytical perspectives, they do so using fundamentally different model architectures and training foundations.

4.1 Agent Architecture and Independence

Each AI agent is implemented as an independent analytical module with its own distinct orientation, evaluation criteria, and underlying language model. The agents do not communicate with each other during the evaluation process, ensuring that each perspective is generated independently. This independence is a deliberate design choice intended to maximize the diversity of the candidate set and prevent the emergence of consensus bias, a well-documented phenomenon in ensemble systems where individual models converge toward similar outputs through implicit information sharing (Dietterich, 2000).

The current architecture employs four primary Signal Agents, one Orchestrator, and a two-level Portfolio Critic. The selection of distinct language models for each role reflects a conviction that model architectural diversity—not just prompt diversity—is essential for genuine analytical pluralism.

4.2 Signal Agent 1: Growth and Momentum (Claude Sonnet 4.5)

The first Signal Agent operates with a growth and momentum analytical bias, implemented through Anthropic's Claude Sonnet 4.5. This agent evaluates securities based on revenue growth trajectories, earnings momentum indicators, relative strength metrics, and forward-looking growth expectations derived from both quantitative data and natural language analysis of corporate communications.

Claude Sonnet 4.5 was selected for this role due to its strong performance in nuanced textual reasoning and its ability to synthesize qualitative and quantitative information—capabilities that are particularly relevant for assessing growth narratives that span financial data, management guidance, and sector dynamics. The agent's output is a ranked assessment of securities exhibiting strong positive momentum characteristics within the investment universe. Securities with declining momentum or negative trajectory indicators are flagged as potential short candidates, providing the construction engine with information relevant to both sides of the signal.

4.3 Signal Agent 2: Risk and Downside Fragility (Mistral Large 2512)

The second Signal Agent focuses on risk-adjusted stability and downside fragility analysis, implemented through Mistral AI's Mistral Large 2512. This agent evaluates securities through the lens of tail risk, drawdown characteristics, volatility regime sensitivity, and balance sheet resilience. Its analytical framework is oriented toward identifying securities that exhibit either unusual resilience or unusual fragility under stress conditions.

Mistral Large 2512 was selected for this role due to its analytical precision in structured quantitative reasoning and its effectiveness in systematic risk decomposition tasks. The agent's fragility assessment extends beyond simple volatility measurement to encompass operating leverage, financial leverage, customer concentration, and competitive positioning factors that influence downside vulnerability. This perspective is particularly valuable in constructing long/short signals, as it provides a natural basis for distinguishing between securities suitable for long positions (those demonstrating structural resilience) and those suitable for short positions (those exhibiting latent fragility that may manifest under adverse conditions).

4.4 Signal Agent 3: Neutral Factor-Based Analysis (Gemini 3 Flash Preview)

The third Signal Agent operates as a neutral, factor-based quantitative analyst, implemented through Google's Gemini 3 Flash Preview. This agent evaluates securities using a systematic factor decomposition approach, assessing exposures to standard risk factors—value, size, quality, low volatility, and profitability (Ang, 2014)—and identifying securities with unusual factor profiles that may represent mispricing or structural dislocations.

Gemini 3 Flash Preview was selected for this role due to its high throughput and efficiency in processing large structured datasets, making it well-suited to the systematic, computation-intensive nature of cross-sectional factor analysis. The agent's perspective is deliberately agnostic with respect to market direction, focusing instead on relative factor characteristics and cross-sectional dispersion within the investment universe. By providing a systematic, theory-grounded perspective, it anchors the candidate set in well-understood risk dimensions and prevents the signal from acquiring unintended factor tilts that might result from the more discretionary orientations of the other agents.

4.5 Signal Agent 4: Contrarian Mean-Reversion (Grok 4)

The fourth Signal Agent operates with a contrarian, mean-reversion bias, implemented through xAI's Grok 4. This agent evaluates securities that have experienced significant price dislocations relative to fundamental anchors, seeking situations where market prices have overreacted to short-term information and are likely to revert toward intrinsic value over the agent's analytical horizon.

Grok 4 was selected for this role due to its distinctive reasoning approach and its willingness to challenge consensus assessments—characteristics that align naturally with a contrarian analytical orientation. The agent's framework incorporates valuation metrics, sentiment indicators, historical mean-reversion patterns, and measures of positioning extremes. This perspective serves as a natural counterbalance to the momentum agent, ensuring that the candidate set includes both trend-following and mean-reverting opportunities. The tension

between momentum and mean-reversion perspectives is a well-known feature of financial markets (Asness, Moskowitz, & Pedersen, 2013), and maintaining both orientations within the analytical layer prevents the signal from becoming systematically biased toward either regime.

4.6 Orchestrator: Final Arbitration (Claude Haiku 4.5)

The outputs of the four Signal Agents are consolidated through an orchestration layer implemented via Anthropic's Claude Haiku 4.5. The Orchestrator performs final arbitration before passing the candidate set to the deterministic construction engine. It does not override individual agent assessments; rather, it resolves conflicts where agents provide contradictory assessments of the same security, flags consensus positions where multiple agents agree on a security's directional assessment, and ensures that the combined candidate set satisfies the minimum diversity requirements specified by the governance framework.

Claude Haiku 4.5 was selected for this role due to its speed, consistency, and low-variance output characteristics—properties that are essential for a coordination function that must operate reliably and transparently. The Orchestrator operates as a lightweight arbitration function, not as a meta-model that reweights agent outputs based on historical accuracy. This design choice reflects the architecture's commitment to transparency: the Orchestrator's logic is fully deterministic and auditable, avoiding the introduction of an additional opaque optimization layer between the agents and the construction engine.

4.7 Portfolio Critic: Two-Level Validation

Before the consolidated candidate set reaches the deterministic construction engine, it passes through a two-level Portfolio Critic that provides an additional validation layer. This critic operates independently from both the Signal Agents and the Orchestrator, serving as a structural safeguard against systematic analytical blind spots.

The first level (L1) is a deterministic Rules Engine that evaluates the candidate set against trend indicators, relative strength index thresholds, sector balance requirements, and beta constraints. This level operates without AI involvement, applying purely mechanical filters that enforce minimum structural quality standards on the candidate set.

The second level (L2) is a shadow analytical layer implemented through Google's Gemini 2.5 Flash, which performs cross-correlation analysis and macro clustering assessment. This shadow layer does not have the authority to modify the candidate set; it generates diagnostic flags that are recorded in the decision log and may trigger governance-level reviews if specific thresholds are breached. The L2 layer serves as an early warning system for structural risks—such as hidden factor concentration or macro regime sensitivity—that may not be visible at the individual security level but emerge at the portfolio level.

5. Deterministic Signal Construction Engine

The deterministic engine represents the core of the architecture. While the AI agents provide analytical richness and candidate diversity, it is the construction engine that transforms these inputs into a fully constrained, deliverable signal. The engine's design reflects a fundamental conviction: in institutional signal generation, the discipline of construction is at least as important as the quality of ideas.

5.1 Rule Hierarchy and Execution Logic

The engine enforces a hierarchical set of rules that are applied sequentially during each construction cycle. The rule hierarchy is organized by priority, with hard constraints (which must never be violated) taking precedence over soft constraints (which represent preferences that may be relaxed if necessary to satisfy higher-priority rules). This hierarchical structure ensures that the most critical governance requirements are always satisfied, even at the expense of less critical optimization objectives.

The primary rules governing signal construction include: equal position sizing across both long and short books; long/short balance to ensure approximately symmetric market exposure; sector diversification constraints limiting concentration in any single sector; correlation limits restricting the maximum pairwise and average correlation between signal positions; liquidity thresholds requiring minimum daily trading volume for all constituents; earnings proximity constraints excluding securities with imminent earnings announcements from new signal inclusion; and turnover limits constraining the percentage of the signal that can be modified in a single construction cycle.

5.2 Equal-Weight Design Philosophy

The choice of equal position sizing deserves particular attention, as it represents a deliberate departure from the optimization-based weighting schemes common in quantitative portfolio management. Traditional mean-variance optimization requires accurate estimates of expected returns and the covariance matrix—parameters that are notoriously difficult to estimate reliably and are subject to significant estimation error, particularly in the expected return dimension (Meucci, 2005).

Estimation errors in these parameters frequently produce extreme, concentrated portfolios that are highly sensitive to small changes in inputs—a phenomenon known as "error maximization" in the portfolio optimization literature. Equal weighting eliminates this source of instability entirely. By assigning identical weights to all positions, the engine removes the need for return forecasts at the signal construction level. This does not mean that all securities are considered equally attractive; the AI agents have already performed analytical differentiation during the candidate selection phase. Rather, equal weighting ensures that the signal's risk characteristics are determined by the governance constraints rather than by an optimization algorithm's sensitivity to parameter estimates.

5.3 Correlation and Diversification Controls

The engine enforces strict correlation limits between signal positions. During each construction cycle, the engine evaluates the pairwise correlation matrix of candidate securities and rejects combinations that would introduce excessive co-movement into the signal. This

constraint operates at multiple levels: within the long book, within the short book, and between long and short positions.

The correlation control mechanism serves several purposes. It reduces the signal's sensitivity to common factor movements that might affect multiple positions simultaneously. It ensures that the signal's return dispersion is not artificially compressed by correlated positions. And it limits the potential for coordinated drawdowns, where multiple positions move adversely at the same time due to shared risk factor exposures (Roncalli, 2013). The specific correlation thresholds are defined in the governance framework and can be adjusted based on the prevailing market correlation regime.

5.4 Event Risk Management

A distinctive feature of the construction engine is its integrated event risk management capability. The engine maintains awareness of scheduled corporate events—principally earnings announcements, but potentially extending to other material events such as FDA decisions, regulatory hearings, or significant product launches—and adjusts constituent eligibility accordingly.

Securities within a defined proximity window of an earnings announcement are excluded from new signal inclusion. This constraint reflects the empirical observation that security returns around earnings announcements are dominated by event-specific information that is fundamentally unpredictable, regardless of the sophistication of the analytical framework. Earnings surprises, guidance revisions, and management commentary can produce single-day moves of five to fifteen percent in individual securities—magnitudes that can overwhelm the analytical assessment embedded in any position. By excluding these high-variance periods from the signal construction process, the engine reduces the probability of outsized position-level losses driven by event-specific outcomes.

5.5 Reproducibility and Immutability

Every execution of the construction engine produces a complete record of the decision process: the candidate set received from the AI agents, the Portfolio Critic's L1 and L2 diagnostic outputs, the sequential application of each governance rule, the securities accepted and rejected at each stage, and the final signal composition with all associated risk metrics. This record is cryptographically hashed using SHA-256, creating an immutable audit trail that cannot be altered after the fact without invalidating the hash.

The combination of deterministic logic and cryptographic immutability ensures that the signal construction process satisfies the most stringent institutional requirements for reproducibility and auditability. Any party with access to the original inputs and governance rules can independently verify the signal that should have been produced, providing a level of transparency that is uncommon in AI-driven investment systems and that addresses one of the primary barriers to institutional adoption.

6. Beta-Neutral Signal Framework

The architecture described in the preceding sections is applied to the generation of beta-neutral long/short signals within the NASDAQ-100 universe. This application context provides a concrete illustration of the framework's capabilities while addressing a signal construction challenge of genuine institutional relevance.

6.1 Investment Universe Selection

The NASDAQ-100 universe is selected as the primary investment universe for several compelling reasons. First, it comprises highly liquid large-capitalization securities, ensuring that signal constituents can be acted upon by subscribers without significant market impact—a critical requirement for institutional-grade signals. Second, the constituent securities span multiple sectors within the technology-heavy index, providing sufficient diversity for sector-balanced long/short construction. Third, the universe is well-covered by sell-side research and corporate disclosure, providing rich input data for the AI analytical agents. Fourth, the index's composition is transparent and rules-based, with clear reconstitution procedures that can be anticipated and incorporated into the signal generation process.

The choice of a defined, liquid universe also supports the architecture's governance objectives. Because the universe is known in advance and its constituents are well-characterized, the governance framework can be calibrated specifically to the statistical properties of the investment universe—correlation structures, sector compositions, liquidity profiles, and event calendars—rather than requiring generic constraints that may be poorly suited to the actual securities in the signal.

6.2 Long/Short Structure

The signal is structured as a balanced long/short composition with equal-weight positions on both sides. The long book comprises securities that the analytical agents have identified as possessing favorable characteristics—strong momentum, fundamental resilience, attractive factor profiles, or mean-reversion potential—subject to the governance constraints imposed by the construction engine.

The short book comprises securities identified as exhibiting the opposite characteristics: deteriorating momentum, fundamental fragility, unattractive factor profiles, or overextension beyond fundamental anchors. The symmetric construction of the long and short books is a deliberate design choice intended to minimize the signal's net market exposure and ensure that the return generation process is driven by security selection rather than market direction.

6.3 Beta Neutrality and Market Exposure Management

The objective of beta neutrality is to construct a signal whose implied returns are substantially independent of the overall direction of the equity market. In a perfectly beta-neutral signal, the systematic component—the portion attributable to broad market movements—would be zero, leaving only the idiosyncratic component that reflects security-specific outcomes (Jacobs & Levy, 1993).

In practice, perfect beta neutrality is neither achievable nor necessarily desirable. Individual security betas vary over time, and the relationship between security returns and market

returns is subject to regime-dependent non-linearities that cannot be fully captured by a static beta estimate. The construction engine therefore targets approximate beta neutrality by balancing the aggregate beta exposure of the long book against the aggregate beta exposure of the short book, subject to a tolerance band defined in the governance rules.

This approach reduces dependence on overall market trends and focuses the signal on the relative performance dispersion within the investment universe. The economic premise is that the AI agents can identify meaningful differences in the relative attractiveness of securities within the NASDAQ-100, even when the aggregate direction of the index is difficult to predict. By removing the market beta component, the architecture isolates the value added by security selection—the domain where AI analytical capabilities are most likely to provide a durable edge.

6.4 Periodic Signal Generation

The signal is regenerated at periodic intervals defined by the governance framework. Each cycle involves a complete re-evaluation of the investment universe by the AI analytical agents, validation by the Portfolio Critic, and a fresh execution of the deterministic construction engine. The regeneration frequency represents a trade-off between the desire to incorporate new information promptly and the turnover implications for subscribers, including transaction costs, market impact, and tax consequences.

The turnover constraint within the construction engine provides an additional mechanism for managing the pace of signal change. Even if the AI agents' assessments shift significantly between cycles, the engine's turnover limit ensures that signal transitions are gradual and controlled, preventing abrupt repositioning that could generate excessive costs for subscribers or signal instability in the investment process.

7. Auditability and Decision Integrity

A central design objective of the system is the creation of comprehensive, immutable decision records that enable full reconstruction of the signal generation process at any point in time. This capability is not an afterthought or an add-on feature; it is a foundational architectural requirement that shapes every aspect of the system's design and represents one of the architecture's most significant differentiators relative to conventional AI-driven investment systems.

7.1 Decision Record Architecture

Each signal construction cycle generates a complete decision record comprising several components. The first component is a snapshot of the candidate securities proposed by each AI Signal Agent, including the agent's internal ranking and the analytical rationale for each candidate. The second component is a record of the Orchestrator's consolidation process, documenting how individual agent outputs were combined into a unified candidate set and how conflicts between agents were resolved. The third component is the Portfolio Critic's L1 deterministic filter results and L2 shadow diagnostic flags. The fourth component is the detailed rule-based filtering sequence applied by the deterministic construction engine, showing which candidates were accepted or rejected at each stage and which specific governance rule was responsible for each decision. The fifth component is the final signal composition, including position sizes, sector allocations, beta exposures, correlation characteristics, and all associated risk metrics.

7.2 Cryptographic Immutability

The complete decision record is cryptographically hashed using SHA-256, creating a unique digital fingerprint that is permanently associated with the signal construction cycle. This hash is stored in an immutable log that provides tamper-evident protection against after-the-fact modification of the decision record. Any alteration to the decision record, however minor, would produce a different hash value, immediately revealing the tampering.

The practical significance of cryptographic immutability extends beyond technical security. It provides institutional subscribers with a verifiable guarantee that the decision process they are reviewing is identical to the process that was actually executed. This guarantee is particularly valuable in regulatory examinations, investor due diligence processes, and internal compliance reviews, where the integrity of the decision record is a fundamental requirement.

7.3 Transparency for Institutional Stakeholders

The combination of deterministic construction logic and immutable decision records creates a level of process transparency that is unusual in AI-driven investment systems. Compliance officers at subscribing institutions can verify that governance rules were correctly applied in every construction cycle. Risk managers can trace the source of specific signal exposures back through the construction process to the original agent assessments. Portfolio managers can understand why specific securities were included or excluded from the signal. External auditors and regulators can independently reconstruct the signal construction process given the original inputs and governance rules.

This transparency is not merely a desirable feature; it may be a necessary condition for the institutional adoption of AI-assisted investment processes. Regulatory frameworks across major jurisdictions—including the European Union's AI Act (European Commission, 2021), evolving SEC guidance on algorithmic trading, and the Bank of England's supervisory expectations for AI in financial services (Bank of England & Financial Conduct Authority, 2019)—increasingly emphasize the importance of explainability and auditability in algorithmic decision-making. Investment managers deploying AI-derived signals must be prepared to demonstrate that their processes satisfy these requirements, and the governance-first architecture provides a natural foundation for doing so.

8. Implications for Institutional AI Deployment

The architecture presented in this paper suggests a broader perspective on the role of artificial intelligence in institutional asset management. Rather than viewing AI as a replacement for traditional portfolio governance—a perspective that has generated considerable anxiety among institutional investors and regulators alike—the framework proposes that AI may be most effective when embedded within deterministic structural frameworks that preserve institutional control and process integrity.

8.1 The Complementarity Thesis

The central thesis of this paper can be stated as follows: AI and deterministic governance are complementary rather than substitutive. AI systems excel at exploring large information spaces, identifying non-obvious patterns, and generating diverse analytical perspectives. Deterministic governance frameworks excel at enforcing discipline, ensuring consistency, and maintaining transparency. Attempting to combine these functions within a single system—for example, by training an AI model to simultaneously optimize for alpha and respect risk constraints—compromises both capabilities and produces systems that are neither analytically powerful nor institutionally robust.

By separating the analytical function (AI Signal Agents) from the construction function (deterministic engine), the proposed architecture allows each component to operate at its maximum effectiveness. The AI agents can explore freely within their analytical domains without concern for signal construction constraints. The deterministic engine can enforce governance rules mechanically without concern for the source or quality of the candidate securities. This separation of concerns is a fundamental principle of robust system design and is directly applicable to the challenge of integrating AI within institutional investment processes.

8.2 Implications for Regulatory Frameworks

The governance-first architecture has significant implications for the evolving regulatory landscape around AI in financial services. Regulators across major jurisdictions have expressed concern about the opacity of AI-driven investment processes and the difficulty of conducting effective oversight of systems whose decision logic is embedded in complex neural network architectures that resist human interpretation (Financial Stability Board, 2017).

The proposed architecture directly addresses these concerns by providing a clear decomposition of the signal generation process into auditable components. The AI agents can be evaluated independently for analytical soundness through standard model validation procedures. The governance rules can be reviewed and approved by compliance functions using existing regulatory frameworks. The deterministic engine can be tested for correct implementation of the approved rules through deterministic verification. And the immutable decision records provide a permanent, verifiable audit trail that supports regulatory examination at any level of detail.

8.3 The Signal-as-a-Service Model

The Signal-as-a-Service delivery model has implications that extend beyond the technical architecture. By separating signal generation from execution, the platform creates a clean

division of regulatory responsibility. The signal provider operates within the domain of research and analytics, not asset management. Subscribing institutions retain full control over execution, custody, and portfolio implementation, integrating the signal within their existing operational and compliance infrastructure.

This model is particularly well-suited to the current regulatory environment, where the boundaries between research, advisory, and management functions are subject to increasing scrutiny. A signal that is delivered as a structured data product—with full audit trails and transparent construction logic—occupies a clearly defined regulatory position that is distinct from discretionary asset management.

8.4 Scalability Across Investment Strategies

While this paper focuses on beta-neutral equity signals within the NASDAQ-100 universe, the architectural principles are potentially applicable to a broader range of investment strategies and asset classes. The fundamental design pattern—AI analytical agents generating candidates within a deterministic construction framework, delivered as a structured signal—can be adapted to different investment universes, different risk management priorities, and different return objectives without fundamental redesign.

For example, the framework could be adapted to credit markets, where AI agents analyze issuer fundamentals and market technicals while a deterministic engine enforces credit quality, duration, and concentration constraints. It could be applied to multi-asset allocation, where AI agents assess regime dynamics and relative value across asset classes while a deterministic engine maintains strategic allocation targets. The modular nature of the architecture supports this kind of generalization while preserving the governance properties that make it suitable for institutional deployment.

8.5 The Human Element

It is important to emphasize that the proposed architecture does not eliminate the human element from the investment process. Humans design the governance rules, select the language models for each analytical role, define the construction constraints, and make strategic decisions about the overall signal approach. The system automates the execution of these decisions but does not replace the judgment that shapes them.

This division of labor reflects a pragmatic view of the comparative advantages of human and artificial intelligence in signal generation. Humans excel at strategic judgment, institutional knowledge, stakeholder management, and the kind of contextual reasoning that is essential for setting appropriate governance parameters. AI systems excel at consistent, systematic analysis across large information sets, rapid processing of diverse data sources, and the maintenance of analytical discipline across many securities simultaneously. The proposed architecture assigns each to the domain where it is most effective, creating a human-AI collaboration that leverages the strengths of both.

9. Limitations and Future Research

This paper focuses on the conceptual architecture of the proposed framework rather than on empirical performance validation. Several important limitations and open research questions merit explicit acknowledgment.

9.1 Absence of Empirical Evidence

The most significant limitation of this paper is the absence of empirical performance results. The architecture is presented as a design framework, not as a validated investment strategy. Future work must include rigorous backtesting across multiple market regimes, out-of-sample validation on data not used in system development, and ideally live signal performance data accumulated over a meaningful time period to assess the practical viability of the approach.

It should be noted, however, that the architecture's primary value proposition is not unconditional alpha generation but rather the combination of AI-assisted analysis with institutional-grade governance. Performance evaluation should therefore consider risk-adjusted returns, process stability, governance compliance, and the consistency of the decision process alongside absolute return measures. A system that generates moderate but consistent risk-adjusted signals within a fully transparent and auditable framework may be more valuable to institutional subscribers than one that generates higher but less predictable returns within an opaque process.

9.2 Agent Design and Model Selection

The specification of AI agent analytical orientations and the selection of specific language models for each role involve design choices that may significantly influence signal characteristics. The current four-agent design—Claude Sonnet 4.5 for growth/momentum, Mistral Large 2512 for risk/fragility, Gemini 3 Flash Preview for neutral/factor-based analysis, and Grok 4 for contrarian/mean-reversion—represents one possible configuration among many. Alternative configurations might include different model assignments, additional agents focused on macroeconomic sensitivity or event-driven analysis, or different combinations of commercial and open-source language models. The optimal number of agents, their analytical orientations, and the most effective model-to-role mapping remain open empirical questions.

9.3 Regime Dependence

The performance characteristics of beta-neutral strategies are known to be regime-dependent. During periods of extreme market stress, correlations across securities tend to increase sharply, potentially reducing the effectiveness of the diversification controls built into the construction engine and compressing the performance dispersion that the signal relies upon. Conversely, during periods of low volatility and high dispersion, the signal may be particularly well-suited to capture returns from security selection. Future research should examine the architecture's behavior across different market regimes and assess whether adaptive modifications to the governance framework can improve risk-adjusted performance across a wider range of market conditions.

9.4 Integration with Macroeconomic Signals

The current architecture focuses on security-level analysis within the investment universe without explicit integration of macroeconomic indicators. Macroeconomic conditions—interest rate environments, credit cycles, fiscal policy changes, geopolitical developments—can have significant effects on the relative attractiveness of different investment styles and factor exposures. Future research may explore the integration of macroeconomic signals, potentially through an additional analytical agent with a macro-focused orientation or through regime-conditional modifications to the governance rules that adjust signal characteristics based on the prevailing macroeconomic environment.

9.5 Transaction Cost Modeling

While the system operates as a signal provider rather than an execution platform, the practical value of the signal to subscribers depends in part on the net-of-cost performance achievable when implementing it. The current design incorporates turnover constraints as a mechanism for managing signal transitions, but does not explicitly model the transaction costs that subscribers would face. Future research should assess the signal's implementability across different capital base sizes and execution capabilities, providing subscribers with guidance on the relationship between signal performance, implementation costs, and optimal rebalancing frequency.

10. Conclusion

Artificial intelligence offers powerful tools for exploring financial markets, but institutional adoption requires systems that respect strict governance principles. The tension between AI's analytical power and institutional requirements for determinism, transparency, and auditability has been a persistent obstacle to the broader deployment of AI within asset management. This paper has presented a framework designed to resolve this tension.

The proposed architecture separates the signal generation process into distinct, auditable domains: an AI-driven analytical layer that generates diverse, multi-perspective candidate assessments through four specialized language models, a two-level Portfolio Critic that provides independent validation, and a deterministic construction engine that transforms candidates into fully constrained, institutionally compliant signals. By subordinating AI-generated assessments to a rule-based construction engine, the system preserves the analytical benefits of multi-agent AI while ensuring that signal formation satisfies institutional requirements for governance, reproducibility, and auditability.

The key architectural innovations can be summarized as follows. The governance-first hierarchy ensures that risk constraints and signal governance rules take precedence over AI-generated analytical assessments, reversing the order commonly found in AI-driven investment systems. The multi-agent analytical layer employs diverse, independent AI perspectives—each implemented through a distinct large language model selected for its alignment with a specific analytical orientation—to generate a rich candidate set. The deterministic construction engine transforms candidate securities into constrained signals through a fully reproducible, rule-based process. Cryptographic immutability through SHA-256 hashing creates tamper-evident decision records that support the most stringent institutional audit requirements. And the Signal-as-a-Service delivery model cleanly separates signal generation from trade execution, preserving governance integrity while allowing institutional subscribers to integrate the signal within their own risk and execution frameworks.

The application of this architecture to beta-neutral long/short signals within the NASDAQ-100 universe demonstrates its practical applicability to a strategy of genuine institutional relevance. However, the architectural principles—separation of analytical and construction functions, deterministic governance, multi-agent diversity, immutable auditability, and API-first signal delivery—are potentially generalizable to a broader range of investment strategies and asset classes.

This separation between idea generation and risk governance may represent a promising path for the integration of artificial intelligence within institutional asset management. Rather than asking whether AI can replace traditional investment processes, the more productive question may be how AI can be embedded within institutional frameworks in ways that enhance analytical capability without compromising governance integrity. The governance-first architecture proposed in this paper offers one answer to that question.

The evolution of AI in financial markets is still in its early stages, and the optimal architecture for institutional deployment remains an open question. We believe that governance-first designs, in which AI serves as an analytical tool within a deterministic structural framework, represent a conceptually sound and practically viable approach to this challenge—one that

respects the hard-won institutional wisdom about the primacy of risk governance while embracing the genuine analytical advances that artificial intelligence can provide.

Author Information

Research Series
AI Market Neutral
Working Paper No. 01

Ivan Alessandro Ongaro is the founder of AI Market Neutral, a research initiative focused on deterministic AI-assisted portfolio construction architectures.

More information: https://aimarketneutral.com

References

Portfolio Construction and Risk Management

Grinold, R. C., & Kahn, R. N. (1999). Active Portfolio Management: A Quantitative Approach for Producing Superior Returns and Controlling Risk (2nd ed.). New York, NY: McGraw-Hill.

Meucci, A. (2005). Risk and Asset Allocation. Berlin, Germany: Springer (Springer Finance).

Roncalli, T. (2013). Introduction to Risk Parity and Budgeting. Boca Raton, FL: Chapman and Hall/CRC.

Factor Investing and Beta-Neutral Strategies

Ang, A. (2014). Asset Management: A Systematic Approach to Factor Investing. New York, NY: Oxford University Press.

Asness, C. S., Moskowitz, T. J., & Pedersen, L. H. (2013). Value and momentum everywhere. Journal of Finance, 68(3), 929–985.

Jacobs, B. I., & Levy, K. N. (1993). Long/short equity investing. Journal of Portfolio Management, 20(1), 52–63.

Machine Learning and AI in Finance

López de Prado, M. (2018). Advances in Financial Machine Learning. Hoboken, NJ: John Wiley & Sons.

Dixon, M., Halperin, I., & Bilokon, P. (2020). Machine Learning in Finance: From Theory to Practice. Cham, Switzerland: Springer.

Rasekhschaffe, K. C., & Jones, R. C. (2019). Machine learning for stock selection. Financial Analysts Journal, 75(3), 70–88.

Governance, Explainability, and Institutional AI

European Commission. (2021). Proposal for a Regulation of the European Parliament and of the Council laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) and amending certain Union legislative acts (COM(2021) 206 final, 2021/0106(COD)). Brussels: European Commission.

Financial Stability Board. (2017). Artificial Intelligence and Machine Learning in Financial Services: Market Developments and Financial Stability Implications. Basel: Financial Stability Board.

Bank of England & Financial Conduct Authority. (2019). Machine learning in UK financial services. London: Bank of England and Financial Conduct Authority.

Multi-Agent Systems and Ensemble Methods

Wooldridge, M. (2009). An Introduction to MultiAgent Systems (2nd ed.). Chichester, UK: John Wiley & Sons.

Dietterich, T. G. (2000). Ensemble methods in machine learning. In J. Kittler & F. Roli (Eds.), Multiple Classifier Systems (Lecture Notes in Computer Science, Vol. 1857, pp. 1–15). Berlin, Germany: Springer.

Market Microstructure and Signal Decay

Hasbrouck, J. (2007). Empirical Market Microstructure: The Institutions, Economics, and Econometrics of Securities Trading. New York, NY: Oxford University Press.

McLean, R. D., & Pontiff, J. (2016). Does academic research destroy stock return predictability? Journal of Finance, 71(1), 5–32.

Disclaimer

This paper is published for discussion and informational purposes only. It does not constitute investment advice, a solicitation to buy or sell securities, or an offer to provide investment management services. The architecture described herein is a conceptual framework for signal generation and does not involve direct management of capital or execution of trades. Past conceptual analysis is not indicative of future performance.

The author may have financial interests in entities related to the subject matter of this paper. Readers should conduct their own independent analysis and consult qualified financial, legal, and regulatory advisors before making investment decisions.

© 2026 Ivan Alessandro Ongaro / IOOX Srl. All rights reserved.